

MicroFuge: A Middleware Approach to Providing Performance Isolation in Cloud Storage Systems

Akshay Singh, **Xu Cui**, Benjamin Cassell, Bernard Wong and Khuzaima Daudjee



June 25, 2014

Storage Resources in Cloud Datacenters

- ▶ Cloud computing allows sharing of resource at the cost of reduced isolation.
- ▶ Storage systems are highly sensitive to performance interference.
- ▶ Lack of performance isolation → Unpredictable latencies.

A Cloud Scenario

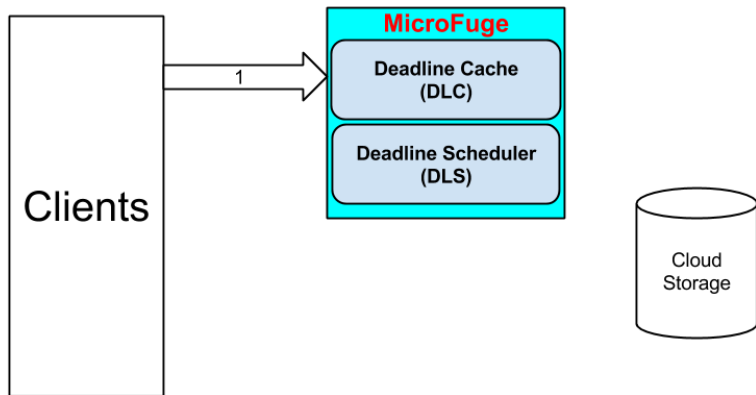
- ▶ In worst case, a particular HTTP request may require 35 database lookups.¹
 - ▶ Response time can add up quickly.
 - ▶ Amazon reported 100ms of latency cost them 1% in sales.²
 - ▶ Google found an extra .5 seconds delay caused 20% drop in search traffic.²
-
- [1] Nathan Farrington and Alexey Andreyev, Facebook's Data Center Network Architecture.
 - [2] Greg Lindem, Make Data Useful, <http://www.scribd.com/doc/4970486/Make-Data-Useful-by-Greg-Linden-Amazon-com>.

Performance Isolation

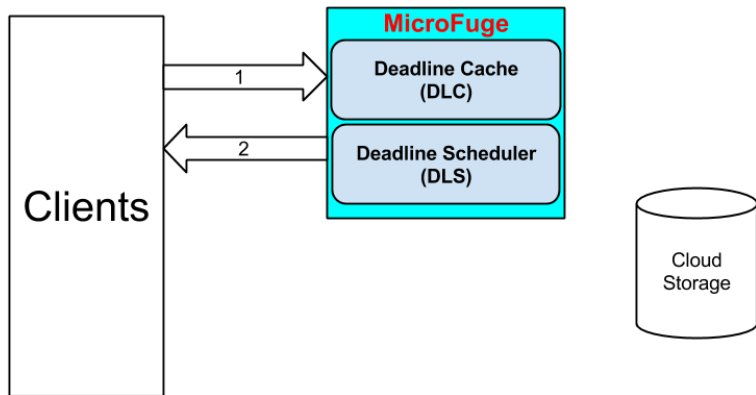
- ▶ Clients want to have performance guarantees in the shared environment.
- ▶ Possible solutions to performance isolation.
 - ▶ Dedicated resources.
 - ▶ Meet clients' response time requirements in the shared environment.
 - ▶ We represent response time requirements with **request deadlines**.
 - ▶ Meeting request deadlines → Performance isolation.

- ▶ A distributed caching and scheduling middleware that provides performance isolation.
 - ▶ **Deadline Cache (DLC)**
 - ▶ Builds a performance model of the system.
 - ▶ Uses multiple LRU queues for deadline-aware eviction.
 - ▶ **Deadline Scheduler (DLS)**
 - ▶ Performs intelligent replica selection.
 - ▶ Implements feedback-driven deadline-aware scheduling.
 - ▶ Optionally performs admission control.
- ▶ Middleware: supports different cloud storage systems.

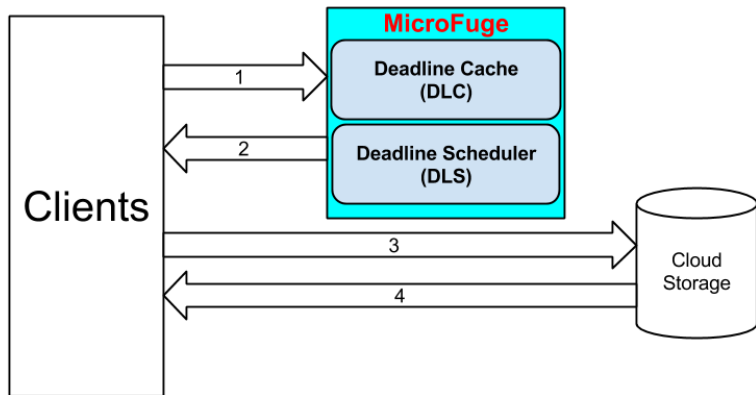
MicroFuge Overview I



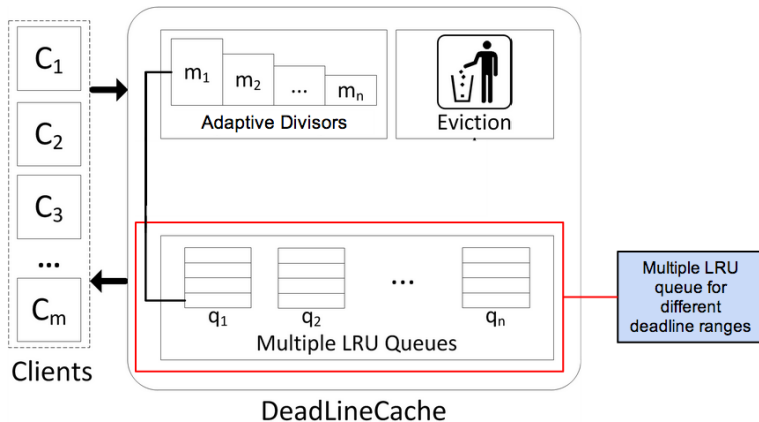
MicroFuge Overview II



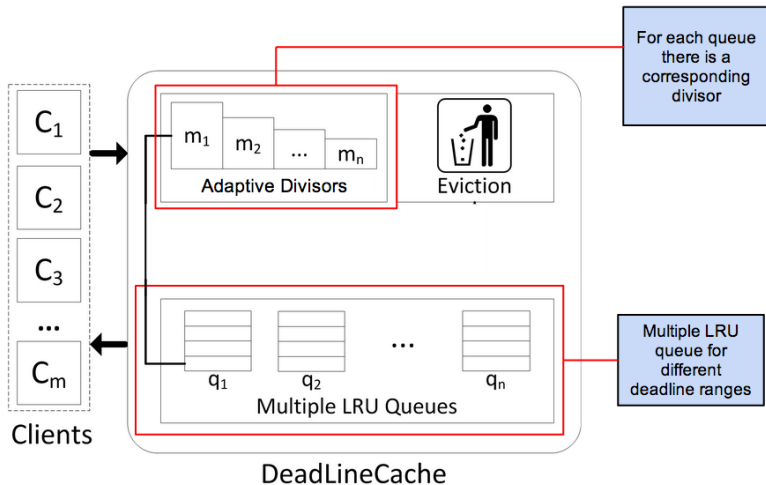
MicroFuge Overview III



Deadline Cache (DLC) - Components



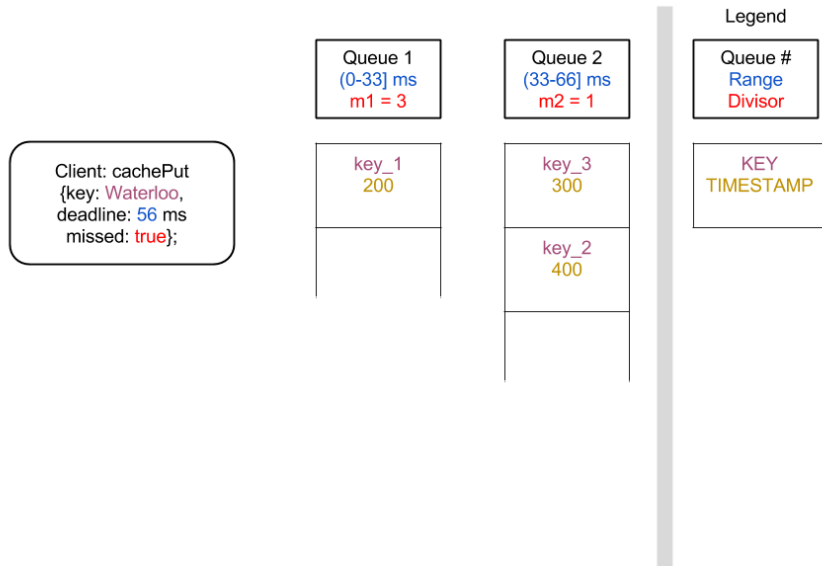
Deadline Cache (DLC) - Components



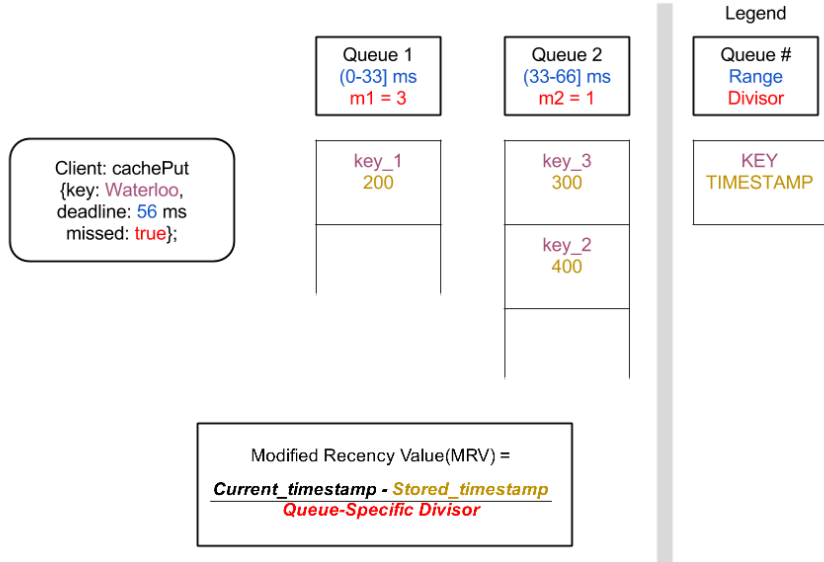
DLC - A Cache Eviction Example (1)

```
Client: cachePut  
{key: Waterloo,  
deadline: 56 ms  
missed: true};
```

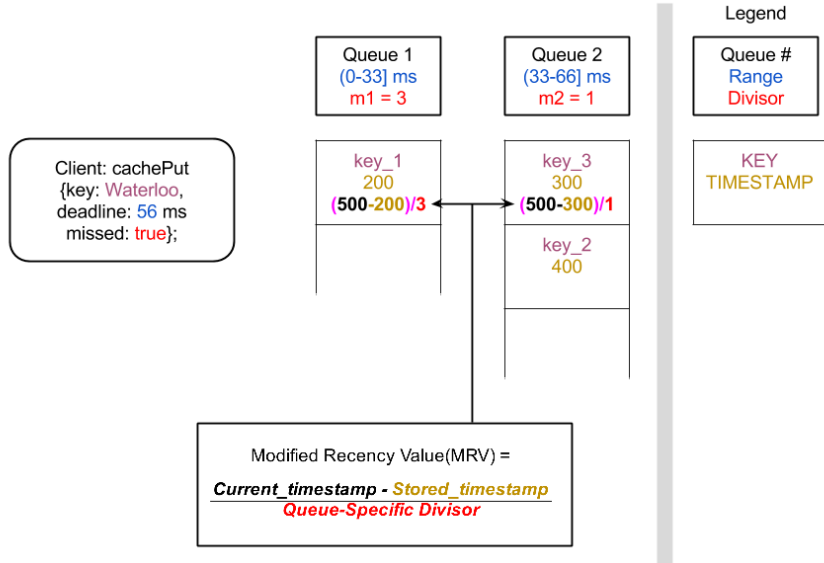
DLC - A Cache Eviction Example (2)



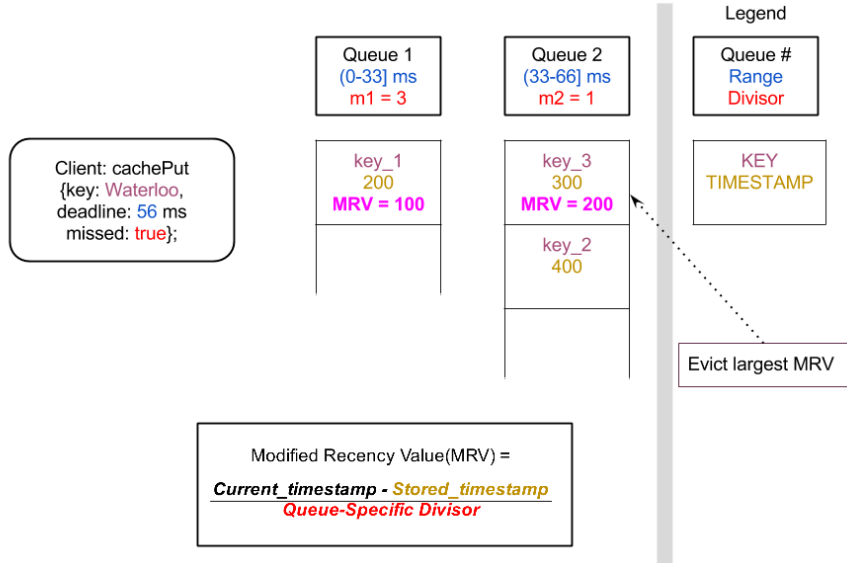
DLC - A Cache Eviction Example (3)




DLC - A Cache Eviction Example (4)



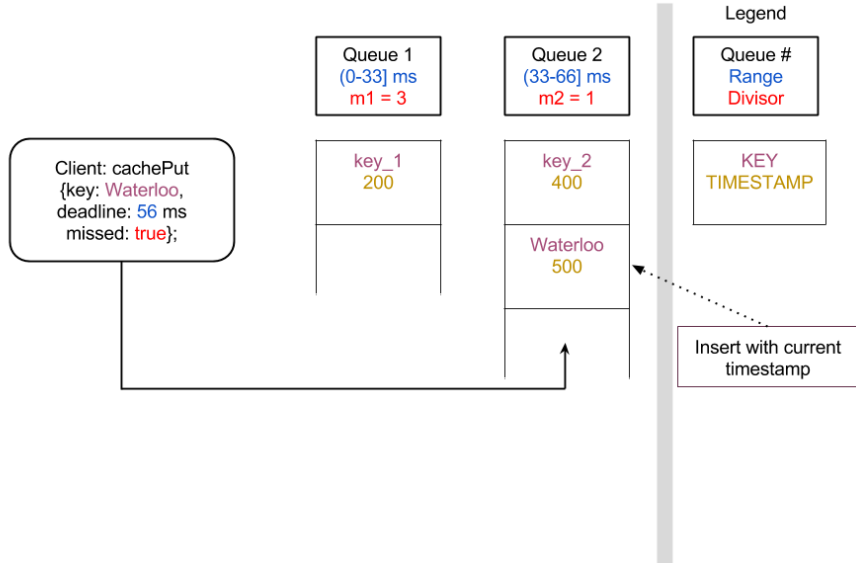
DLC - A Cache Eviction Example (5)



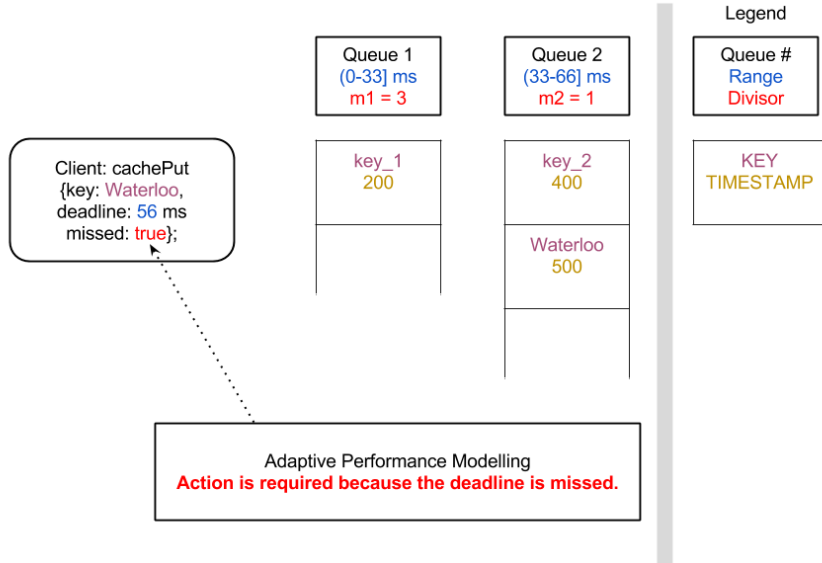
DLC - A Cache Eviction Example (6)



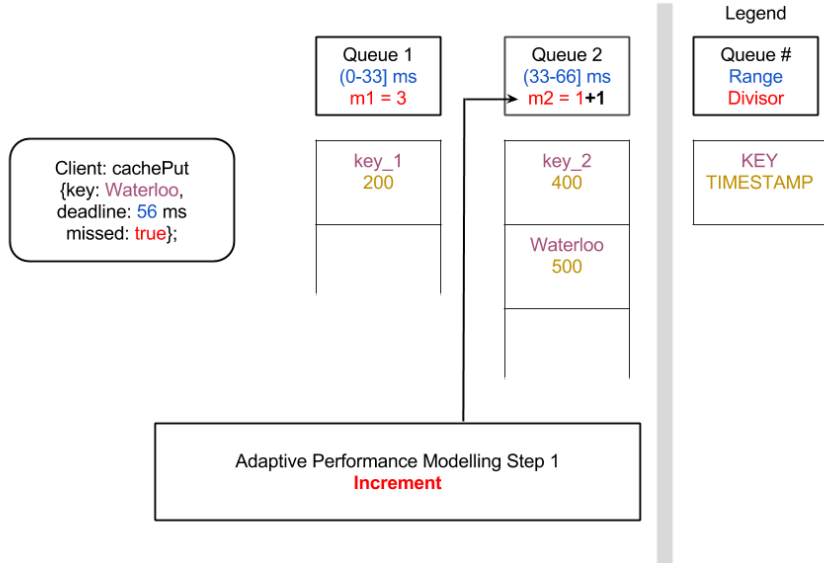
DLC - A Cache Eviction Example (7)



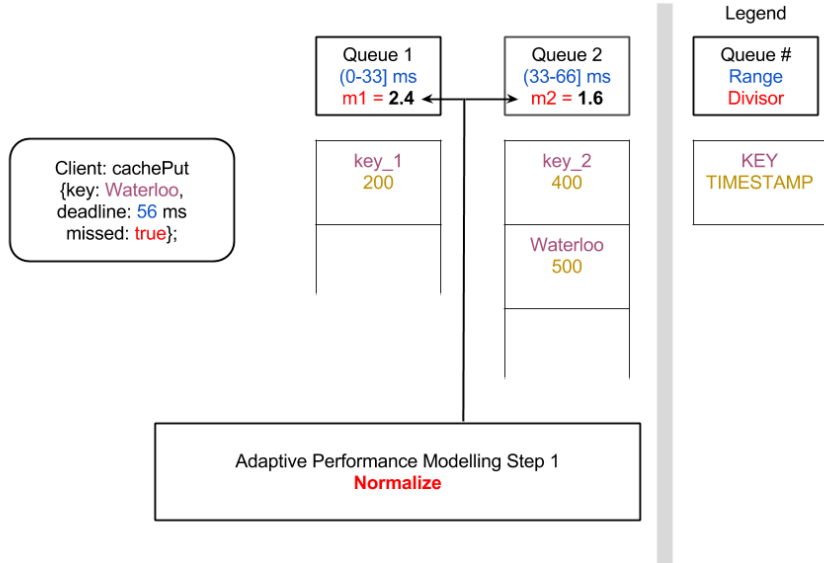
DLC - A Cache Eviction Example (8)



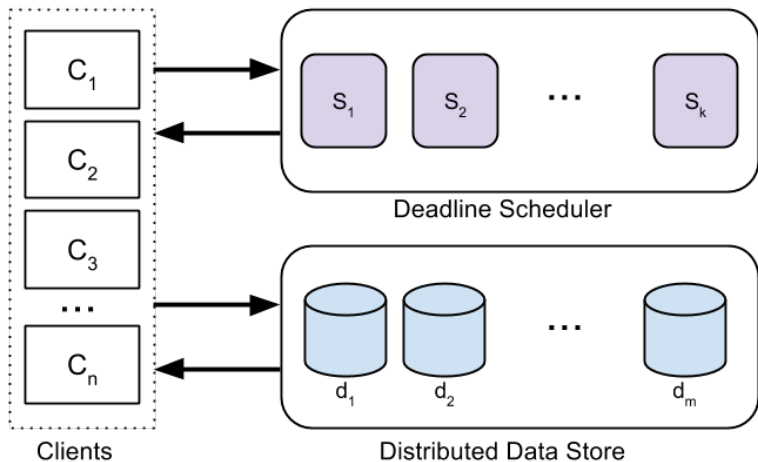
DLC - A Cache Eviction Example (9)



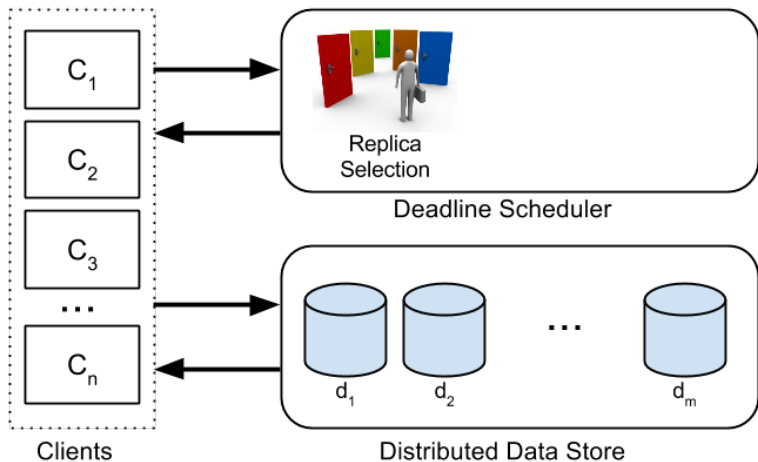
DLC - A Cache Eviction Example (10)



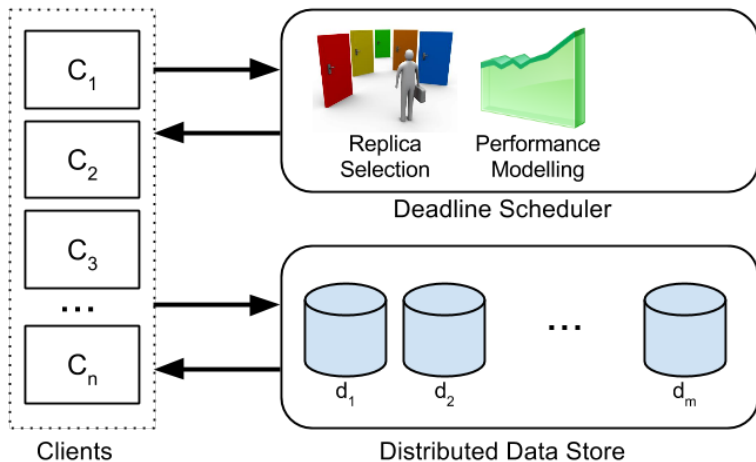
Deadline Scheduler (DLS) High-level Architecture I



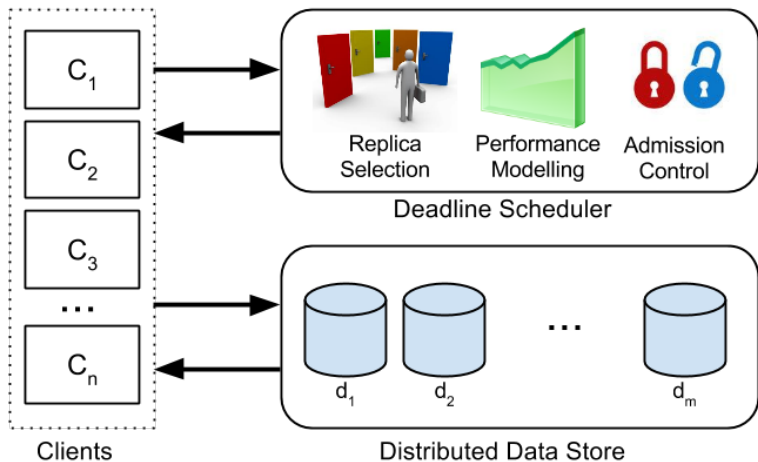
Deadline Scheduler (DLS) High-level Architecture II



Deadline Scheduler (DLS) High-level Architecture III




Deadline Scheduler (DLS) High-level Architecture IV




DLS - An Example (1)

- ▶ The client wants to perform a value lookup for the key Waterloo.




DLS - An Example (2)

- The client begins by issuing a cache lookup to DLC.



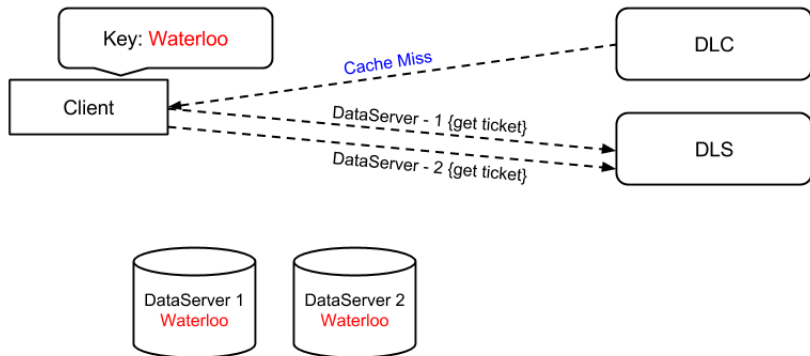
DLS - An Example (3)

- Issue two *get ticket* requests concurrently.



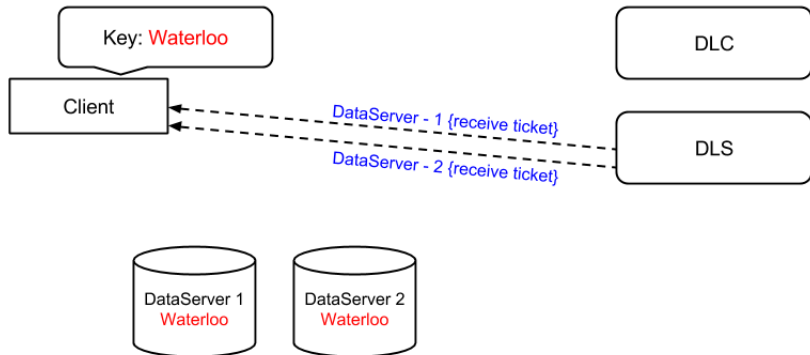
DLS - An Example (4)

- If the item is not in the cache, the client waits for DLS to return the tickets.



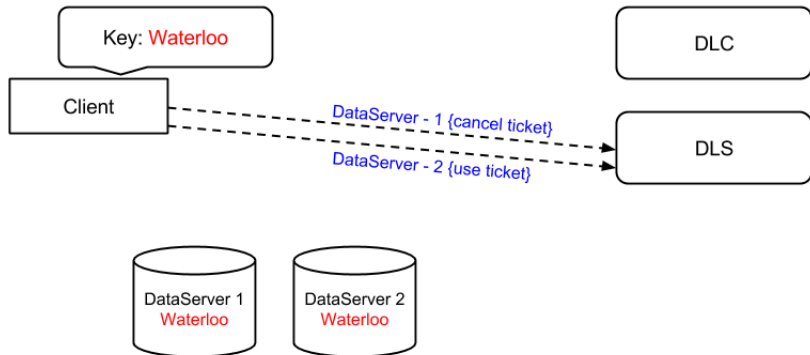
DLS - An Example (5)

- Returned tickets contain extra information to help the client to make an informed decision.





DLS - An Example (6)

- The client makes a blocking call to the selected DLS and waits for its turn to access the data server.




DLS - An Example (7)

- A snapshot of scheduler's pending queue.




DLS - An Example (8)

- The new item is inserted according to earliest deadline first ordering.



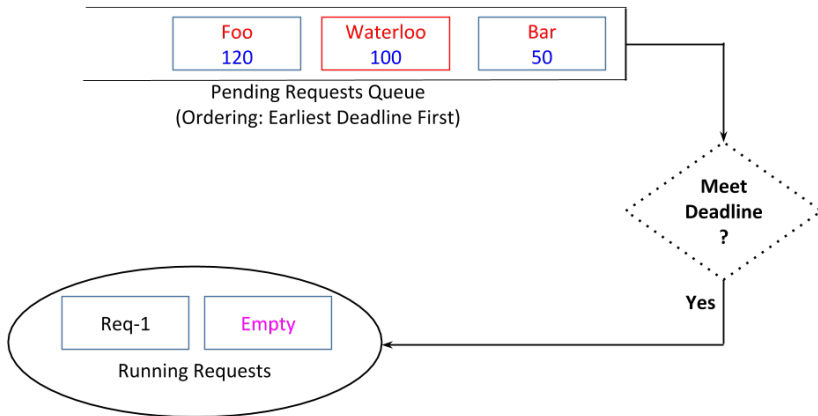
DLS - An Example (9)

- Let's assume one of the running requests just completed.



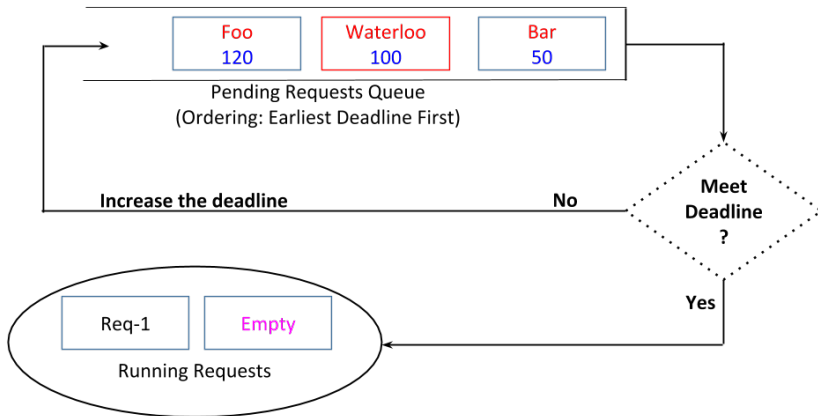
DLS - An Example (10)

- If the request deadline can be met, it will take one of the empty slots inside the running request pool.



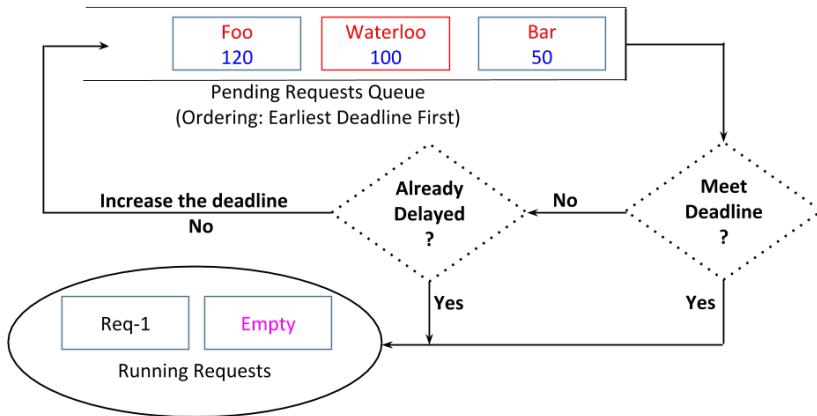
DLS - An Example (11)

- If request deadline cannot be met, DLS may increase the request's deadline and insert the request back into the queue.




DLS - An Example (12)

- ▶ The push-back can happen at most once to prevent starvation.



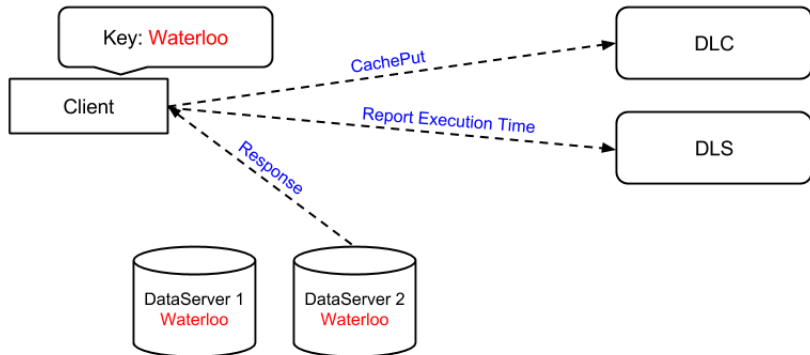
DLS - An Example (14)

- The client issues the read request to the data server.




DLS - An Example (15)

- ▶ After receiving the response, the client reports the execution time and concurrently inserts the data into the cache.



Experimental Setup - The Cluster

- Twenty-node test cluster on AWS. Each cluster node is an m1.medium EC2 instance.



Experimental Setup - Details

- ▶ DataServer - Simple key-value store that uses leveldb.
- ▶ We use a replication factor of 3.
- ▶ Benchmarking System - Modified version of Yahoo! Cloud Serving Benchmark (YCSB).
 - ▶ Assign deadlines to each key.

Range	Percentage
10-30ms	20%
30-100ms	30%
100-1000ms	50%

- ▶ Data Set - 80 million records, 86.4 GB in size.
- ▶ Cache - Total capacity of 19.2GB.

Deadline-Aware Caching - DLC




Figure : Cache hit rate for 192 concurrent clients with DLC and Memcached.

Deadline-Aware Caching - Full MicroFuge




Figure : Cache hit rate for 192 concurrent clients with DLC + DLS and Memcached.

Deadline Miss Rate - DLC




Figure : Deadline miss rate for 192 concurrent clients with DLC and Memcached.

Deadline Miss Rate - Full MicroFuge




Figure : Deadline miss rate for 192 concurrent clients with DLC + DLS and Memcached.

Conclusion

- ▶ Predictable performance is necessary in multi-tenant environments.
- ▶ MicroFuge tackles the performance isolation problem with its deadline-aware caching and scheduling middleware.
- ▶ MicroFuge reduces deadline miss rate from 17.5% to 7.7% and it can be as low as 4.7% if we turn on the admission control.

Thank You.

DLS - Admission Control

- ▶ Bound the fraction of requests that miss their deadlines.
- ▶ Requests are rejected in two situations.
 - ▶ The request will miss its own deadline.
 - ▶ The new request will cause already accepted requests to miss their deadlines.
- ▶ Provides a system parameter β as a knob to control the percentage of deadline misses.

Experimental Results - Deadline Miss with Admission Control




Figure : Deadline miss rate for 192 concurrent clients with DLC + DLS + AC and Memcached.

Experimental Results - Tunable Admission Control




Figure : Deadline miss vs. rejection rates with respect to various values of system parameter β for 192 clients.

MicroFuge at a Glance

- ▶ Middleware for popular key-value storage.
- ▶ A modified version of the CRUD operation interface.

```
// READ interface
public String read(String key, double deadline, boolean bestEffort);

// A sample READ operation with a 15 milliseconds deadline
String myVal = read("myKey", 15, true);
```

Figure : MicroFuge *read* operation interface.